

# NATURAL LANGUAGE PROCESSING

## Textual Insight Engine

Uncovering Literary Patterns with Advanced NLP & Information  
Extraction


Book used for this project:

**The Alchemist by Paulo Coelho**

# Overview

In this project, we will use different Natural Language Processing models to analyze a book. We will apply PoS tagging and create a Bi-gram table using the NLP frameworks defined in Python. The actual code is uploaded on GitHub but the screenshots of the output are uploaded using the code on Jupyter Notebook with some changes in the variable names.

## Book Used



The Alchemist by Paulo Coelho

## Github Link for the code and Implementation

<https://github.com/SwastikDala>

## Problem Statement

Performing the following steps on a book used for the project.

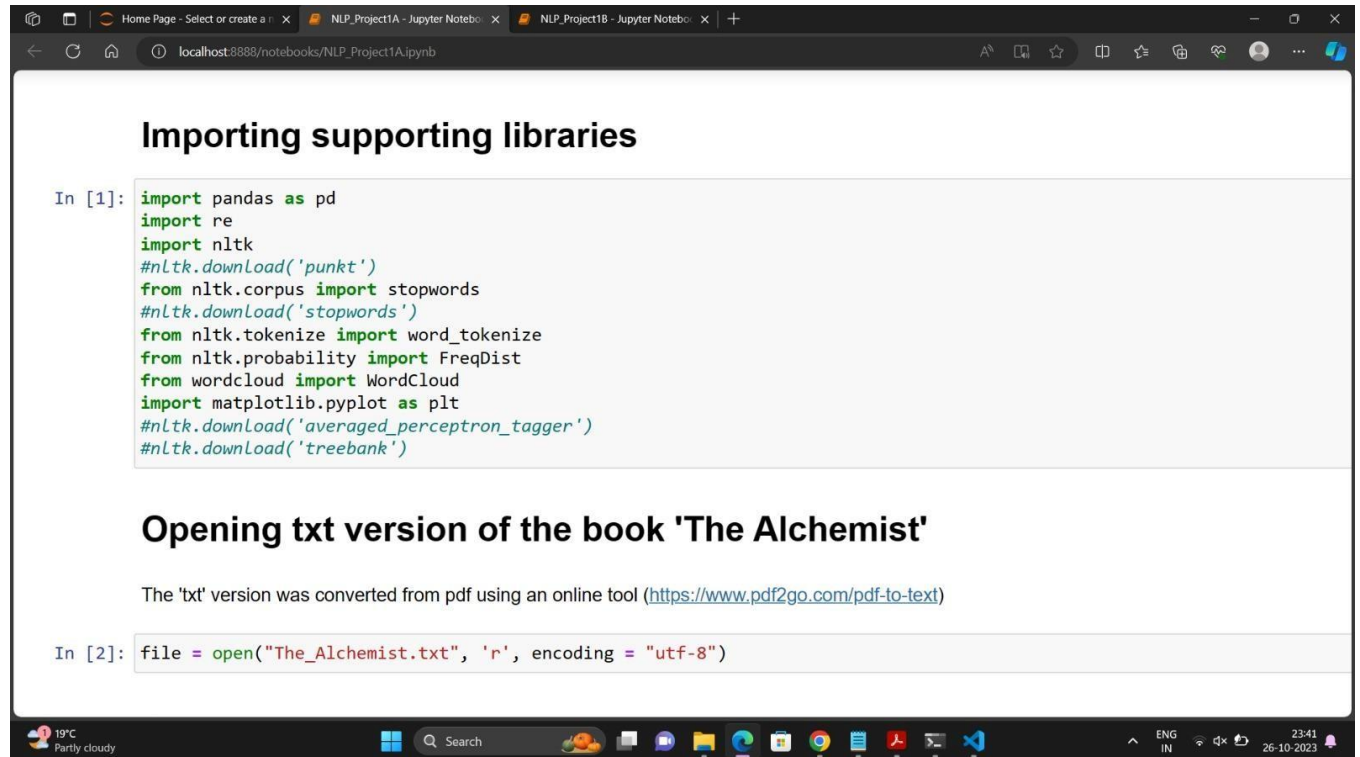
1. Import the text from the book in txt format.
2. Perform Pre-processing on the text.
3. Perform tokenization.
4. Create a frequency distribution table for the text.
5. Create a word cloud for the tokens.
6. Perform PoS tagging.
7. Create a bi-gram probability table for the largest chapter in the book.
8. Play the Shannon game using the bi-gram table.
9. Analyze the accuracy of the bi-gram model.

## Data Description

The pdf book is first transformed into txt format. However, before we can use any of the Python NLP frameworks, we must first eliminate any irrelevant data from the converted text or data that has no impact on the frequency distribution table or the bi-gram model. The content includes photos, the name of the book that appears at the bottom-right of each page, page numbers, chapter names, and a watermark from the website where the book was downloaded. This data may have a negative impact on NLP frameworks and so must be eliminated. The actual application of Python begins in the following section.

# Data Pre-Processing

Before starting data pre-processing, we will first import the book.



The screenshot shows a Jupyter Notebook window with two cells. The first cell is titled "Importing supporting libraries" and contains code to import pandas, re, nltk, and download nltk corpora. The second cell is titled "Opening txt version of the book 'The Alchemist'" and contains code to open a text file. The notebook is running on a local host at localhost:8888.

```
In [1]: import pandas as pd
import re
import nltk
#nltk.download('punkt')
from nltk.corpus import stopwords
#nltk.download('stopwords')
from nltk.tokenize import word_tokenize
from nltk.probability import FreqDist
from wordcloud import WordCloud
import matplotlib.pyplot as plt
#nltk.download('averaged_perceptron_tagger')
#nltk.download('treebank')
```

### Importing supporting libraries

The 'txt' version was converted from pdf using an online tool (<https://www.pdf2go.com/pdf-to-text>)

```
In [2]: file = open("The_Alchemist.txt", 'r', encoding = "utf-8")
```

Data preprocessing is the process of cleaning and organizing the unstructured text data to prepare it for analysis. We will remove all the irrelevant text data as described in the previous section with the help of Regular Expressions.

1. Remove the name of the book.
2. Remove author name.
3. Remove chapter names.
4. Remove page numbers.
5. Remove special characters.

```
Home Page - Select or create a n... NLP_Project1A - Jupyter Notebo... NLP_Project1B - Jupyter Notebo... +
localhost:8888/notebooks/NLP_Project1A.ipynb
In [4]: # Regular expression pattern to match "PAULO COELHO"
author_pattern = r'PAULO COELHO '
# Use 're.sub()' method to replace the matched pattern with an empty string
text = re.sub(author_pattern, '', text)

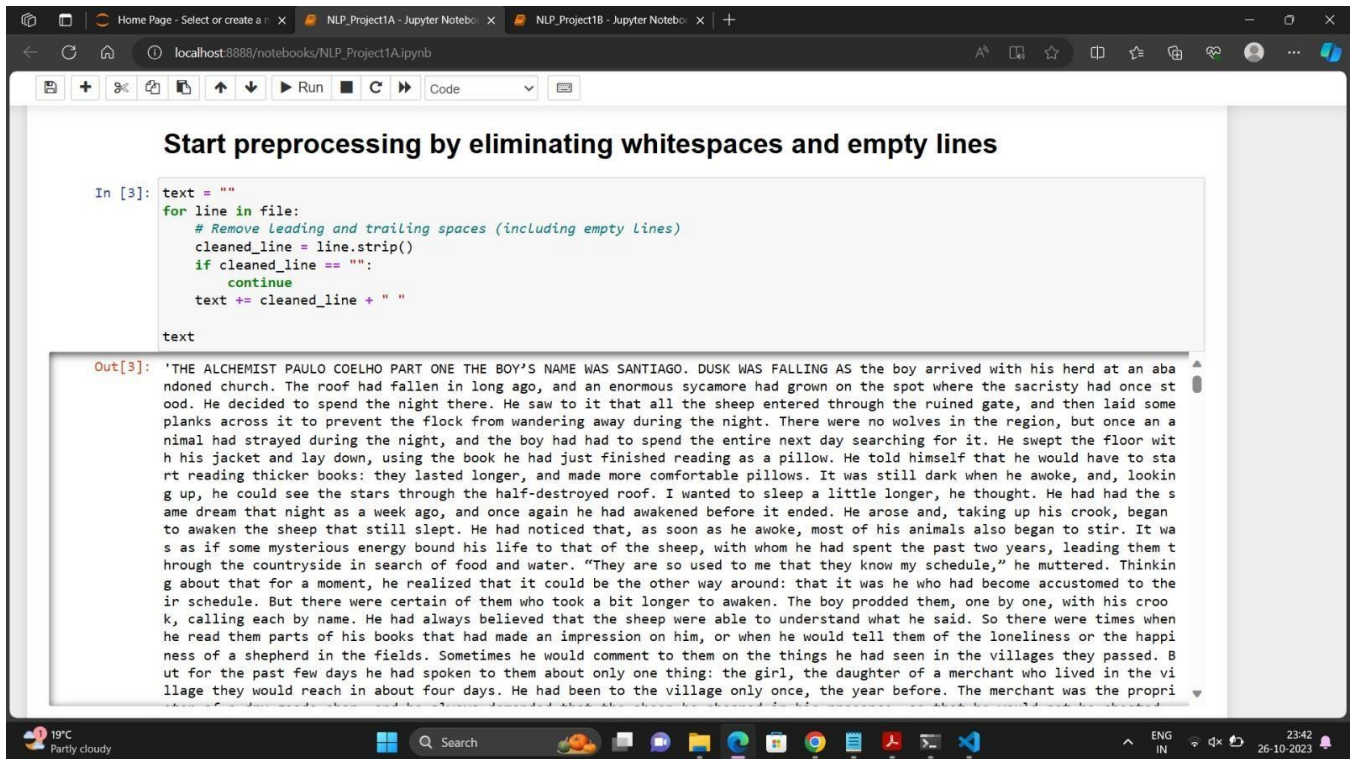
# Regular expression to match "THE ALCHEMIST"
book_pattern = r'THE ALCHEMIST '
'''
# Below code Indicates that only one instance is found in the entire document
matches = re.findall(book_pattern, text)
num_matches = len(matches)
print("Number of matches:", num_matches)
'''
# Replacing pattern with an empty string
text = re.sub(book_pattern, '', text)

# Regular expression to match with chapter number, e.g. "PART ONE", "PART TWO", etc
part_pattern = r'PART [A-Z]+'
# Replacing pattern with an empty string
text = re.sub(part_pattern, '', text)

# Regular expression to match with page number
page_pattern = r'page [0-9]+'
# Replacing pattern with an empty string
text = re.sub(page_pattern, '', text)

# Special character representation using regular expression
specialchar_pattern = r'^a-zA-Z0-9\s'
# Removing special characters
text = re.sub(specialchar_pattern, '', text)
```

## 6. Remove extra whitespace, tabs and newlines.

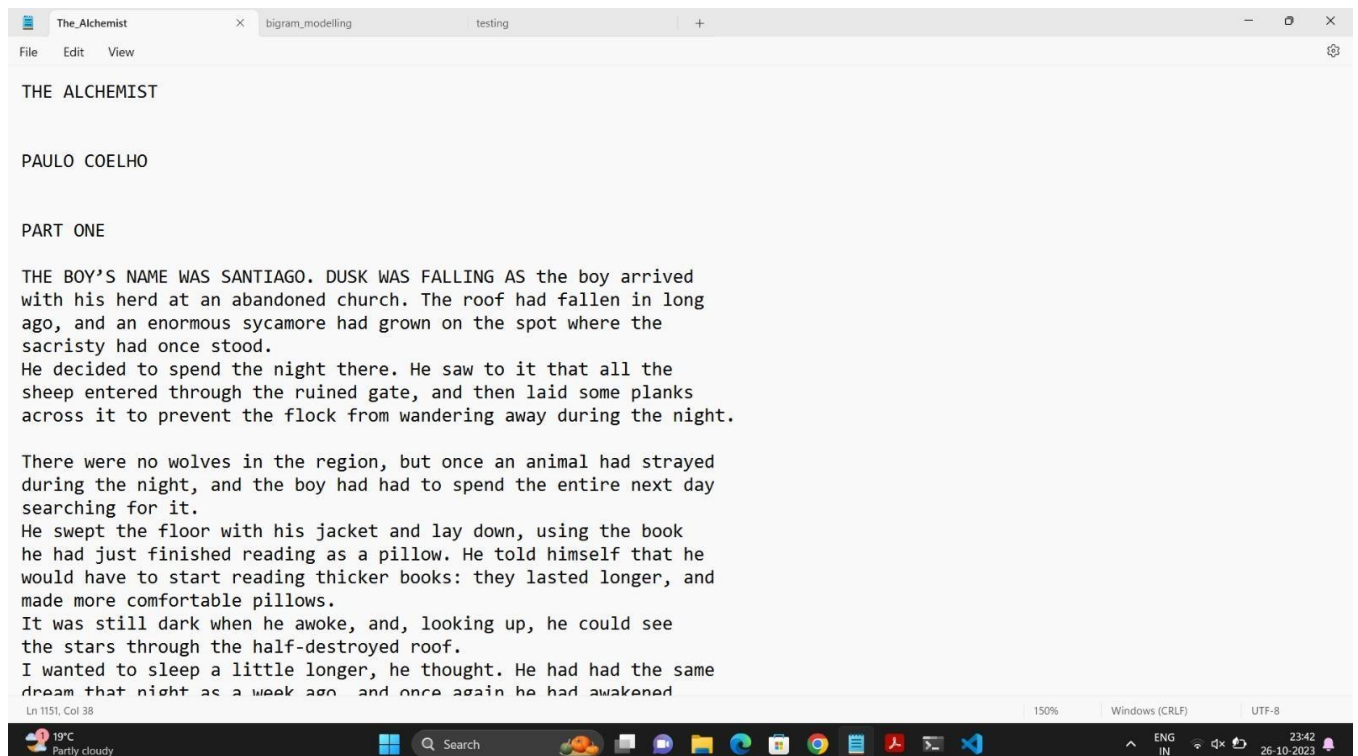


The screenshot shows a Jupyter Notebook interface with a code cell and its output. The code cell contains a Python script that reads a file, strips leading and trailing spaces from each line, and concatenates the cleaned lines into a single string. The output cell shows the resulting text, which is the first chapter of 'The Alchemist' by Paulo Coelho, formatted with proper indentation and line breaks.

```
In [3]: text = ""
for line in file:
    # Remove leading and trailing spaces (including empty lines)
    cleaned_line = line.strip()
    if cleaned_line == "":
        continue
    text += cleaned_line + " "
text
```

```
Out[3]: 'THE ALCHEMIST PAULO COELHO PART ONE THE BOY'S NAME WAS SANTIAGO. DUSK WAS FALLING AS the boy arrived with his herd at an abandoned church. The roof had fallen in long ago, and an enormous sycamore had grown on the spot where the sacristy had once stood. He decided to spend the night there. He saw to it that all the sheep entered through the ruined gate, and then laid some planks across it to prevent the flock from wandering away during the night. There were no wolves in the region, but once an animal had strayed during the night, and the boy had had to spend the entire next day searching for it. He swept the floor with his jacket and lay down, using the book he had just finished reading as a pillow. He told himself that he would have to start reading thicker books: they lasted longer, and made more comfortable pillows. It was still dark when he awoke, and, looking up, he could see the stars through the half-destroyed roof. I wanted to sleep a little longer, he thought. He had had the same dream that night as a week ago, and once again he had awakened before it ended. He arose and, taking up his crook, began to awaken the sheep that still slept. He had noticed that, as soon as he awoke, most of his animals also began to stir. It was as if some mysterious energy bound his life to that of the sheep, with whom he had spent the past two years, leading them through the countryside in search of food and water. "They are so used to me that they know my schedule," he muttered. Thinking about that for a moment, he realized that it could be the other way around: that it was he who had become accustomed to their schedule. But there were certain of them who took a bit longer to awaken. The boy prodded them, one by one, with his crook, calling each by name. He had always believed that the sheep were able to understand what he said. So there were times when he read them parts of his books that had made an impression on him, or when he would tell them of the loneliness or the happiness of a shepherd in the fields. Sometimes he would comment to them on the things he had seen in the villages they passed. But for the past few days he had spoken to them about only one thing: the girl, the daughter of a merchant who lived in the village they would reach in about four days. He had been to the village only once, the year before. The merchant was the propri
```

## Unstructured Data



The screenshot shows a text editor window with the output of the preprocessing script. The text is formatted with proper indentation and line breaks, matching the output shown in the Jupyter Notebook screenshot.

```
THE ALCHEMIST

PAULO COELHO

PART ONE

THE BOY'S NAME WAS SANTIAGO. DUSK WAS FALLING AS the boy arrived
with his herd at an abandoned church. The roof had fallen in long
ago, and an enormous sycamore had grown on the spot where the
sacristy had once stood.
He decided to spend the night there. He saw to it that all the
sheep entered through the ruined gate, and then laid some planks
across it to prevent the flock from wandering away during the night.

There were no wolves in the region, but once an animal had strayed
during the night, and the boy had had to spend the entire next day
searching for it.
He swept the floor with his jacket and lay down, using the book
he had just finished reading as a pillow. He told himself that he
would have to start reading thicker books: they lasted longer, and
made more comfortable pillows.
It was still dark when he awoke, and, looking up, he could see
the stars through the half-destroyed roof.
I wanted to sleep a little longer, he thought. He had had the same
dream that night as a week ago, and once again he had awakened
```

## Output:



## As the next step of tokenisation we convert each word to lowercase

```
In [5]: text = text.lower()  
text
```

```
Out[5]: 'the boys name was santiago dusk was falling as the boy arrived with his herd at an abandoned church the roof had fallen in 1  
ong ago and an enormous sycamore had grown on the spot where the sacristy had once stood he decided to spend the night there  
he saw to it that all the sheep entered through the ruined gate and then laid some planks across it to prevent the flock from  
wandering away during the night there were no wolves in the region but once an animal had strayed during the night and the bo  
y had had to spend the entire next day searching for it he swept the floor with his jacket and lay down using the book he had  
just finished reading as a pillow he told himself that he would have to start reading thicker books they lasted longer and ma  
de more comfortable pillows it was still dark when he awoke and looking up he could see the stars through the halfdestroyed r  
oof i wanted to sleep a little longer he thought he had had the same dream that night as a week ago and once again he had awa  
kened before it ended he arose and taking up his crook began to awaken the sheep that still slept he had noticed that as soon  
as he awoke most of his animals also began to stir it was as if some mysterious energy bound his life to that of the sheep wi  
th whom he had spent the past two years leading them through the countryside in search of food and water they are so used to  
me that they know my schedule he muttered thinking about that for a moment he realized that it could be the other way around  
that it was he who had become accustomed to their schedule but there were certain of them who took a bit longer to awaken the  
boy prodded them one by one with his crook calling each by name he had always believed that the sheep were able to understand  
what he said so there were times when he read them parts of his books that had made an impression on him or when he would tel  
l them of the loneliness or the happiness of a shepherd in the fields sometimes he would comment to them on the things he had  
seen in the villages they passed but for the past few days he had spoken to them about only one thing the girl the daughter o  
f a merchant who lived in the village they would reach in about four days he had been to the village only once the year befor  
e the merchant was the proprietor of a dry goods shop and he always demanded that the sheep be sheared in his presence so tha  
t he would not be cheated. Gerd had told the boy about the sheep and he had taken his sheep there indeed some 11 years ago 1
```

The changed text lacks numbers, special characters (as seen by the fact that the entire text is on the same line), watermarks, and so on.

This concludes the project's data pre-processing phase. The following part depicts the data preparation procedure, which includes text tokenization and stop word removal.

## Data Preparation

This consists of two steps:

- Tokenization

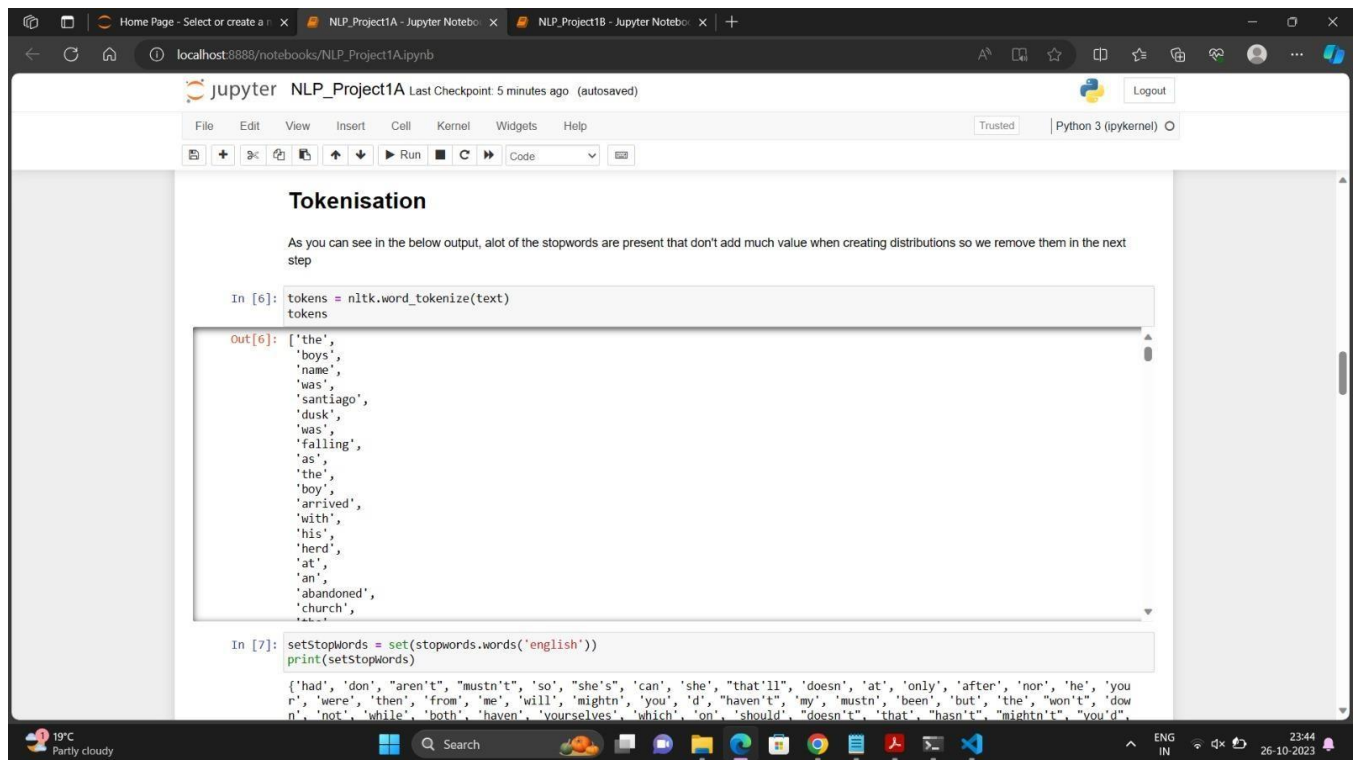
- Removal of stop words

These steps are done using the built-in functions in the NLTK library in Python.

1) Tokenization: It is the process of breaking a stream of textual data into words, terms, sentences, symbols or other meaningful elements called tokens.



2) Removal of stopwords: These are the most common words in any natural language. Since they do not add much value to the document they must be removed before creating a word cloud or a frequency distribution table.



The screenshot shows a Jupyter Notebook interface with the following content:

### Tokenisation

As you can see in the below output, alot of the stopwords are present that don't add much value when creating distributions so we remove them in the next step

```
In [6]: tokens = nltk.word_tokenize(text)
tokens
```

```
Out[6]: ['the',
'boys',
'name',
'was',
'santiago',
'dusk',
'was',
'falling',
'as',
'the',
'boy',
'arrived',
'with',
'his',
'herd',
'at',
'an',
'abandoned',
'church',
'at']
```

```
In [7]: setStopWords = set(stopwords.words('english'))
print(setStopWords)
```

```
{'had', 'don', 'aren't', 'mustn't', 'so', 'she's', 'can', 'she', 'that'll', 'doesn't', 'at', 'only', 'after', 'nor', 'he', 'you r', 'were', 'then', 'from', 'me', 'will', 'mightn', 'you', 'd', 'haven't', 'my', 'mustn', 'been', 'but', 'the', 'won't', 'dow n', 'not', 'while', 'both', 'haven', 'yourselves', 'which', 'on', 'should', 'doesn't', 'that', 'hasn't', 'mightn't', 'you'd', 'is', 'are', 'be', 'it', 'of', 'on', 'the', 'to', 'was', 'were', 'with', 'without', 'you're'}
```

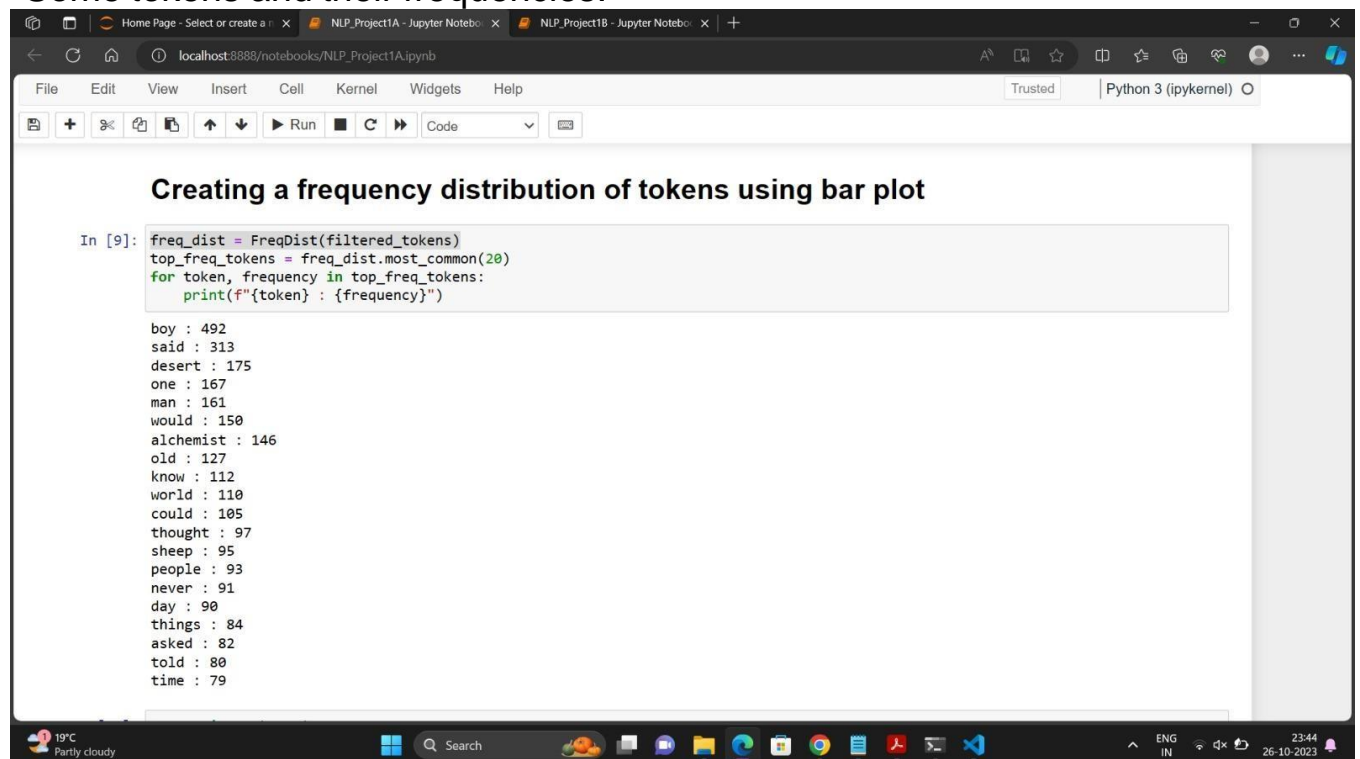
It can be observed that the stopwords like 'is', 'are' and so on are removed from the textual data. Now we can create the frequency distribution table.

# Frequency Distribution Table

This table gives the frequency of each of the tokens generated from the textual data. After tokenization, we create a frequency distribution table directly using a built-in function called `FreqDist(tokens)` in the NLTK library.

We will only store the 20 most common words (tokens), that is, the first 20 tokens when the frequencies of the tokens are arranged in descending order.

Some tokens and their frequencies:



The screenshot shows a Jupyter Notebook interface with a code cell and its output. The code cell is titled "Creating a frequency distribution of tokens using bar plot" and contains the following Python code:

```
In [9]: freq_dist = FreqDist(filtered_tokens)
top_freq_tokens = freq_dist.most_common(20)
for token, frequency in top_freq_tokens:
    print(f"{token} : {frequency}")
```

The output of the code is a list of 20 tokens and their frequencies, displayed as text:

```
boy : 492
said : 313
desert : 175
one : 167
man : 161
would : 150
alchemist : 146
old : 127
know : 112
world : 110
could : 105
thought : 97
sheep : 95
people : 93
never : 91
day : 90
things : 84
asked : 82
told : 80
time : 79
```

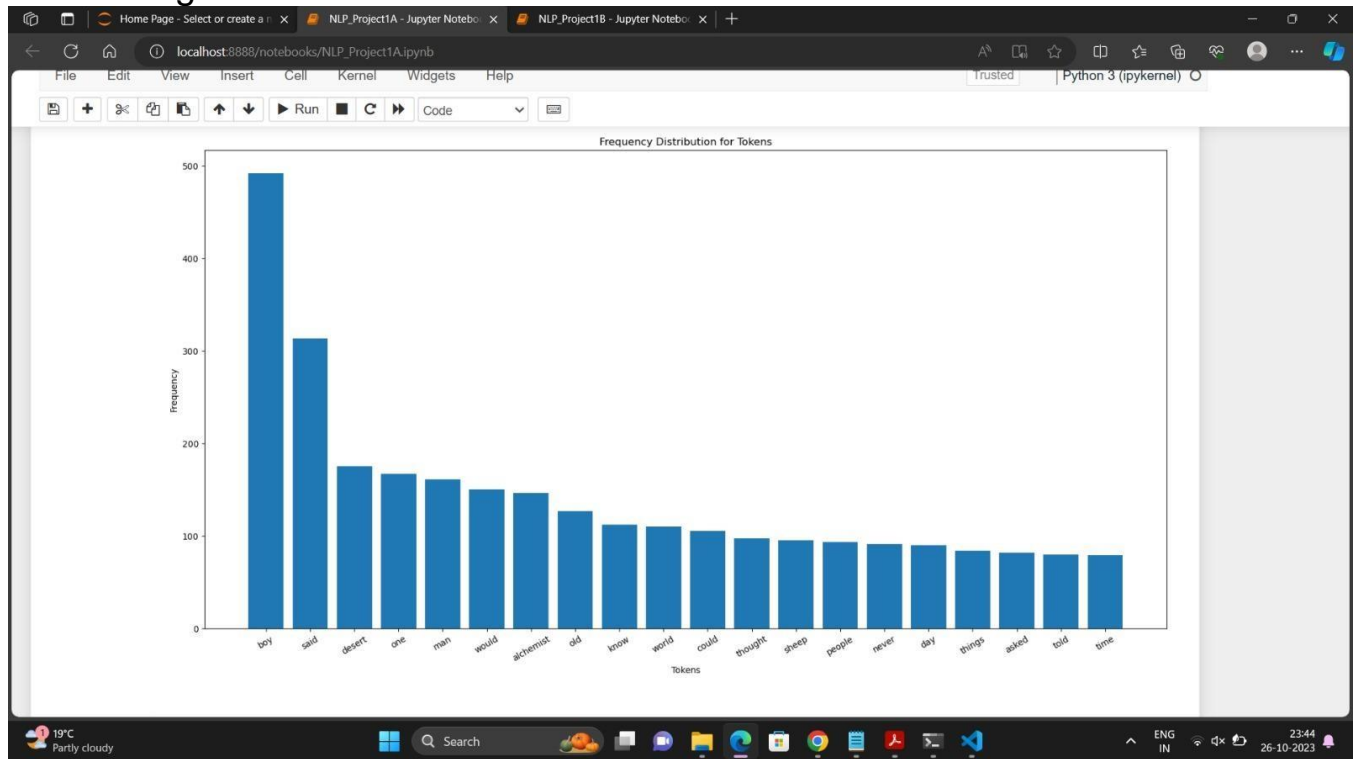
The token 'boy' has the highest frequency which is logical because he is the protagonist of the book.

The frequency distribution table is stored in a non-graphical format and hence must be converted into a more visualizable format. One option is the Word Cloud. It can be created by using WordCloud function using the wordcloud and matplotlib library.

[illegible]

Another option is to create a plot between the words and their frequencies. In this project, we will create a histogram using the 'matplotlib' library.

## The histogram:

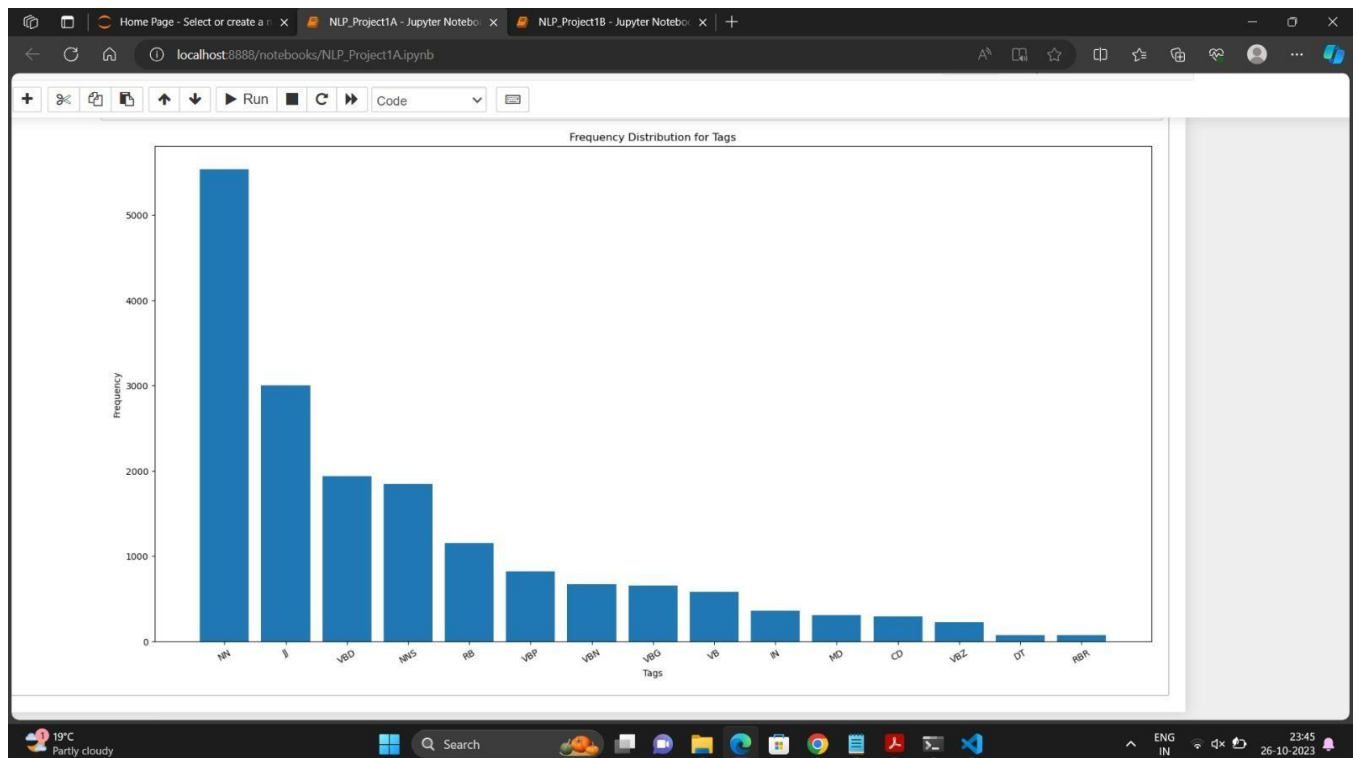


# Part of Speech Tagging

It is a process of assigning each word in the textual data to its corresponding equivalent in part of speech. This includes Nouns, Verbs, Adjectives etc.

This is done using a tag set. In this project, we will use the Penn Treebank Corpus which has 36 PoS tags.

Distribution of some of the tags:



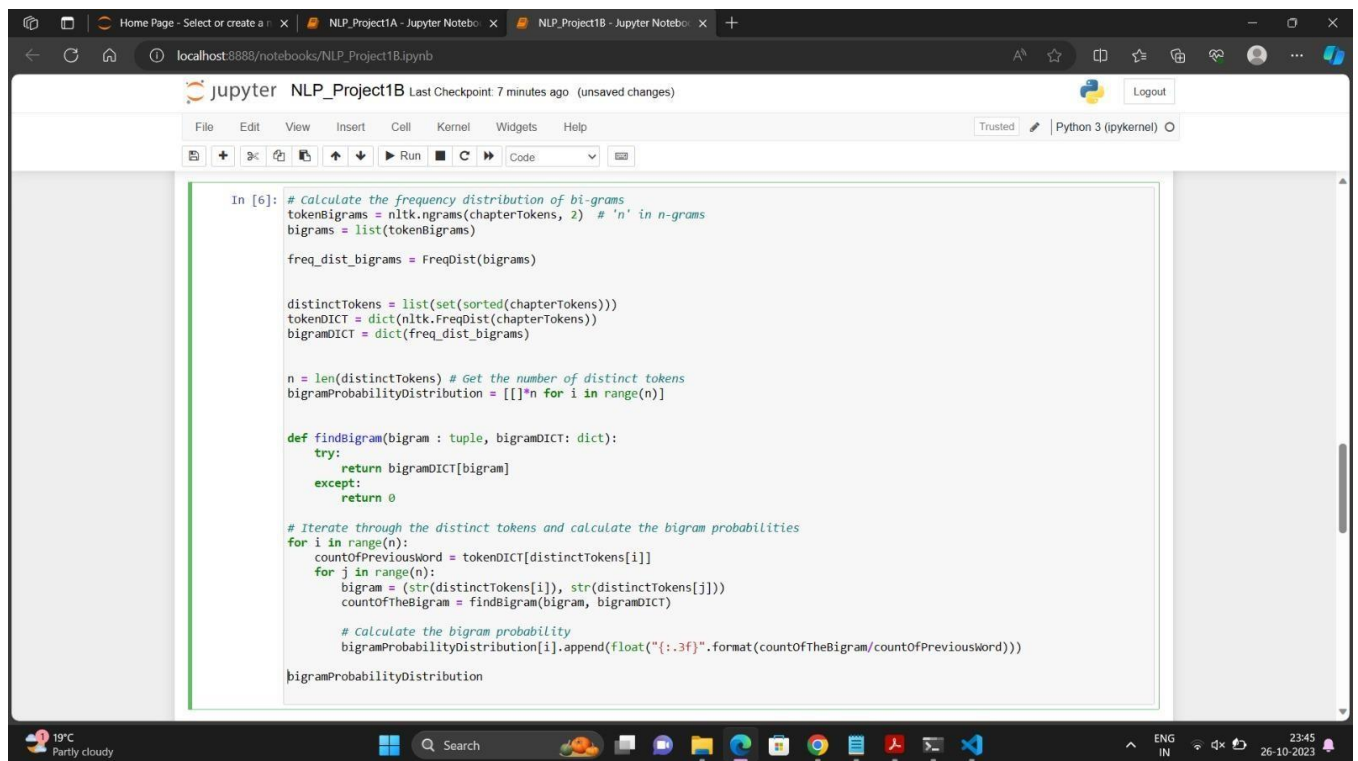


It can be observed that most tokens are tagged as nouns (NN) followed by adjectives (JJ) and so on.

## Bi-gram Model

In this project, we will calculate the bigram probabilities of the words in the largest chapter of the book. The bigram model assumes that the probability of a word appearing in a sentence depends only upon the previous word. To create the bi-gram table, we will first import the largest chapter, pre-process the textual data, tokenize it, and then create the bigram table. First we Import the largest chapter.

Pre-Processing the data and Tokenizing it:



```
In [6]: # Calculate the frequency distribution of bi-grams
tokenBigrams = nltk.ngrams(chapterTokens, 2) # 'n' in n-grams
bigrams = list(tokenBigrams)

freq_dist_bigrams = FreqDist(bigrams)

distinctTokens = list(set(sorted(chapterTokens)))
tokenDICT = dict(nltk.FreqDist(chapterTokens))
bigramDICT = dict(freq_dist_bigrams)

n = len(distinctTokens) # Get the number of distinct tokens
bigramProbabilityDistribution = [[]*n for i in range(n)]

def findBigram(bigram : tuple, bigramDICT: dict):
    try:
        return bigramDICT[bigram]
    except:
        return 0

# Iterate through the distinct tokens and calculate the bigram probabilities
for i in range(n):
    countOfPreviousword = tokenDICT[distinctTokens[i]]
    for j in range(n):
        bigram = (str(distinctTokens[i]), str(distinctTokens[j]))
        countOfTheBigram = findBigram(bigram, bigramDICT)

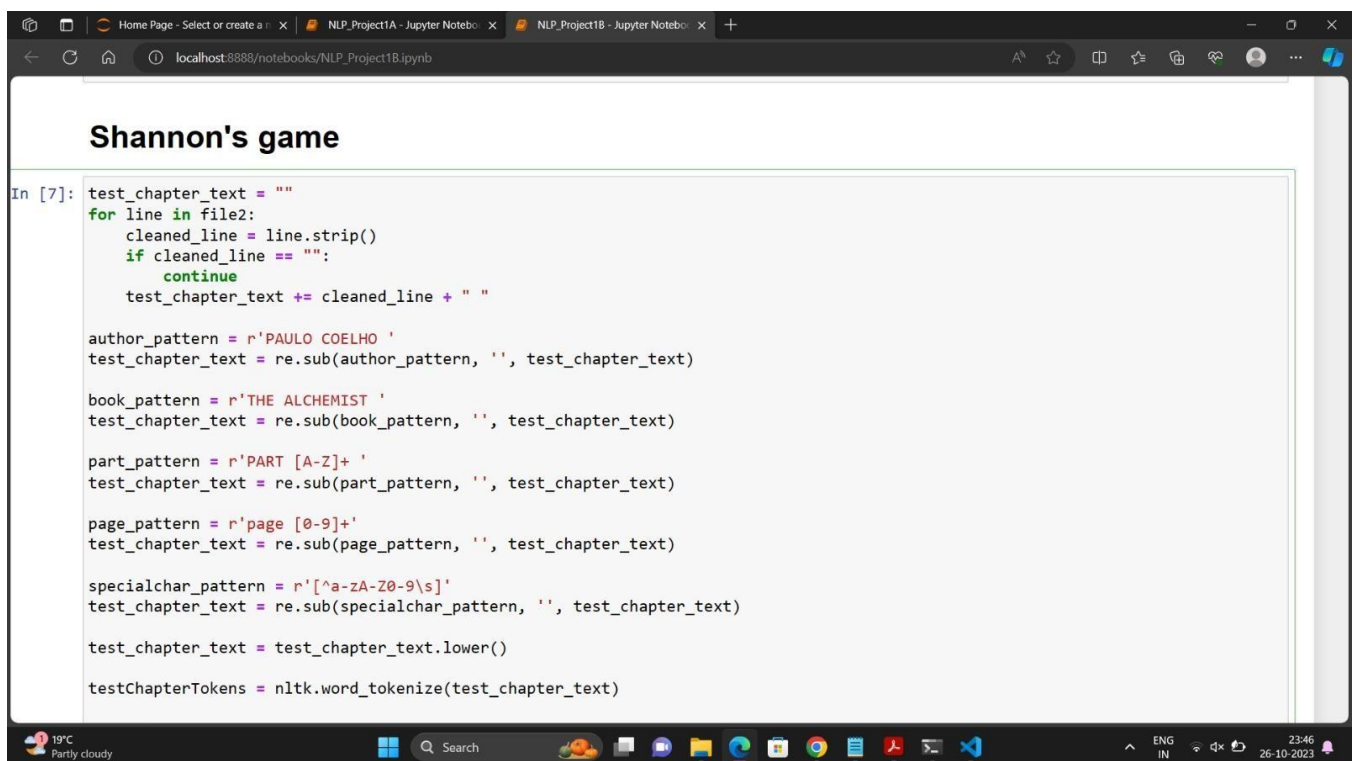
        # Calculate the bigram probability
        bigramProbabilityDistribution[i].append(float("{:.3f}".format(countOfTheBigram/countOfPreviousword)))

bigramProbabilityDistribution
```

# Shannon Game

This game is used to test the accuracy of a word guessing algorithm. Here, using the bi-gram model based upon the largest chapter in the book, we fill in the blanks in another chapter chosen and then calculate the accuracy by comparing the guessed word with the actual word in the chapter.

We first perform the basic steps (importing, pre-processing and tokenization) on a chosen chapter of the book.



```
In [7]: test_chapter_text = ""
for line in file2:
    cleaned_line = line.strip()
    if cleaned_line == "":
        continue
    test_chapter_text += cleaned_line + " "

author_pattern = r'PAULO COELHO '
test_chapter_text = re.sub(author_pattern, '', test_chapter_text)

book_pattern = r'THE ALCHEMIST '
test_chapter_text = re.sub(book_pattern, '', test_chapter_text)

part_pattern = r'PART [A-Z]+'
test_chapter_text = re.sub(part_pattern, '', test_chapter_text)

page_pattern = r'page [0-9]+'
test_chapter_text = re.sub(page_pattern, '', test_chapter_text)

specialchars_pattern = r'^a-zA-Z0-9\s'
test_chapter_text = re.sub(specialchar_pattern, '', test_chapter_text)

test_chapter_text = test_chapter_text.lower()

testChapterTokens = nltk.word_tokenize(test_chapter_text)
```

Now we create 'n' blanks in the text by replacing the tokens created (in the original text) with blanks. Here, we will take  $n=300$ . We will then play the Shannon game.

The blanks are created and then they are filled using the bi-gram probabilities table created using the largest chapter in the book.

Now, we will calculate the accuracy of our model.

The accuracy of the model comes out to be 11.66.

In the project's second phase, various Natural Language Processing techniques will be employed to analyze the book chosen for this project (The Alchemist). The subsequent step involves executing 'Entity Recognition' utilizing Python's NLP frameworks. The code is uploaded on GitHub and the screenshots are from the code on Jupyter Notebook and GitHub.

## **Problem Statement**

We aim to perform the following tasks:

1. Recognize all entities in the book and recognize all entity types.
2. Use different performance measures to evaluate the method used. We will take random passages from the book for this.
3. Repeat the above step three times.
4. Generate a TF-IDF vector for all the chapters in the book.
5. Find similarity between the chapters using similarity measure.
6. Create a gradient table.

## **Libraries Used**

In this round, we will primarily use the 'spaCy' library in Python. spaCy is a free open-source library for Natural Language Processing in Python. It features NER, POS tagging, dependency parsing, word vectors and makes information extraction and general-purpose natural language processing easier.

# Installing spaCy Library

```
In [1]: import pandas as pd
import re
import nltk
#nltk.download('punkt')
from nltk.corpus import stopwords
#nltk.download('stopwords')
from nltk.tokenize import word_tokenize
from nltk.probability import FreqDist
from wordcloud import WordCloud
import matplotlib.pyplot as plt
#nltk.download('averaged_perceptron_tagger')
#nltk.download('treebank')
from nltk.util import ngrams
import random
import sys
import spacy
from spacy import displacy
import numpy as np
import seaborn as sns
from sklearn.metrics import confusion_matrix
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
```

# Data Preprocessing

```
In [3]: text = ""
for line in file:
    # Remove leading and trailing spaces (including empty lines)
    cleaned_line = line.strip()
    if cleaned_line == "":
        continue
    text += cleaned_line + " "
text
```

```
Out[3]: 'THE ALCHEMIST PAULO COELHO PART ONE THE BOY'S NAME WAS SANTIAGO. DUSK WAS FALLING AS the boy arrived with his herd at an abandoned church. The roof had fallen in long ago, and an enormous sycamore had grown on the spot where the sacristy had once stood. He decided to spend the night there. He saw to it that all the sheep entered through the ruined gate, and then laid some planks across it to prevent the flock from wandering away during the night. There were no wolves in the region, but once an animal had strayed during the night, and the boy had had to spend the entire next day searching for it. He swept the floor with his jacket and lay down, using the book he had just finished reading as a pillow. He told himself that he would have to start reading thicker books: they lasted longer, and made more comfortable pillows. It was still dark when he awoke, and, looking up, he could see the stars through the half-destroyed roof. I wanted to sleep a little longer, he thought. He had had the same dream that night as a week ago, and once again he had awakened before it ended. He arose and, taking up his crook, began to awaken the sheep that still slept. He had noticed that, as soon as he awoke, most of his animals also began to stir. It was as if some mysterious energy bound his life to that of the sheep, with whom he had spent the past two years, leading them through the countryside in search of food and water. "They are so used to me that they know my schedule," he muttered. Thinking about that for a moment, he realized that it could be the other way around: that it was he who had become accustomed to their schedule. But there were certain of them who took a bit longer to awaken. The boy prodded them, one by one, with his crook, calling each by name. He had always believed that the sheep were able to understand what he said. So there were times when he read them parts of his books that had made an impression on him, or when he would tell them of the loneliness or the happiness of a shepherd in the fields. Sometimes he would comment to them on the things he had seen in the villages they passed. But for the past few days he had spoken to them about only one thing: the girl, the daughter of a merchant who lived in the village they would reach in about four days. He had been to the village only once, the year before. The merchant was the propri
```

```
In [4]: # Regular expression pattern to match "PAULO COELHO"
author_pattern = r'PAULO COELHO '
# Use 're.sub()' method to replace the matched pattern with an empty string
text = re.sub(author_pattern, '', text)

# Regular expression to match "THE ALCHEMIST"
book_pattern = r'THE ALCHEMIST '
'''
# Below code Indicates that only one instance is found in the entire document
matches = re.findall(book_pattern, text)
num_matches = len(matches)
print("Number of matches:", num_matches)
'''

# Replacing pattern with an empty string
text = re.sub(book_pattern, '', text)

# Regular expression to match with chapter number, e.g. "PART ONE", "PART TWO", etc
part_pattern = r'PART [A-Z]+'
# Replacing pattern with an empty string
text = re.sub(part_pattern, '', text)

# Regular expression to match with page number
page_pattern = r'page [0-9]+'
# Replacing pattern with an empty string
text = re.sub(page_pattern, '', text)

# Special character representation using regular expression
specialchar_pattern = r'[^a-zA-Z0-9\s] '

# Removing special characters
text = re.sub(specialchar_pattern, '', text)
```

```
In [5]: preprocessed_text
```

```
Out[5]: 'The boys name was Santiago dusk was falling as the boy arrived with his herd at an abandoned church The roof had fallen in l
ong ago and an enormous sycamore had grown on the spot where the sacristy had once stood He decided to spend the night there
He saw to it that all the sheep entered through the ruined gate and then laid some planks across it to prevent the flock from
wandering away during the night There were no wolves in the region but once an animal had strayed during the night and the bo
y had had to spend the entire next day searching for it He swept the floor with his jacket and lay down using the book he had
just finished reading as a pillow He told himself that he would have to start reading thicker books they lasted longer and ma
de more comfortable pillows It was still dark when he awoke and looking up he could see the stars through the halfdestroyed r
oof I wanted to sleep a little longer he thought He had had the same dream that night as a week ago and once again he had awa
kened before it ended He arose and taking up his crook began to awaken the sheep that still slept He had noticed that as soon
as he awoke most of his animals also began to stir It was as if some mysterious energy bound his life to that of the sheep wi
th whom he had spent the past two years leading them through the countryside in search of food and water They are so used to
me that they know my schedule he muttered Thinking about that for a moment he realized that it could be the other way around
that it was he who had become accustomed to their schedule But there were certain of them who took a bit longer to awaken The
boy prodded them one by one with his crook calling each by name He had always believed that the sheep were able to understand
what he said So there were times when he read them parts of his books that had made an impression on him or when he would tel
l them of the loneliness or the happiness of a shepherd in the fields Sometimes he would comment to them on the things he had
seen in the villages they passed But for the past few days he had spoken to them about only one thing the girl the daughter o
f a merchant who lived in the village they would reach in about four days He had been to the village only once the year before
e The merchant was the proprietor of a dry goods shop and he always demanded that the sheep be sheared in his presence so tha
t he would not be bothered if he had to tell the shepherd the sheep had to be taken to his house about 1000 to 1500 some 1000
```

## Entity Recognition

To generate all the entities, we will first import the "spaCy" library and use it to process the text.



```
In [6]: nlp = spacy.load("en_core_web_sm")
```

```
In [7]: bookNER = nlp(preprocessed_text)
```

Now iterating through all the entities in the book using the library.

```
In [8]: for ent in bookNER.ents:  
        print(f"Entity: {ent.text}, Type: {ent.label_}")
```

```
Entity: Santiago, Type: PERSON  
Entity: the night, Type: TIME  
Entity: the night, Type: TIME  
Entity: the night, Type: TIME  
Entity: next day, Type: DATE  
Entity: a week ago, Type: DATE  
Entity: the past two years, Type: DATE  
Entity: one, Type: CARDINAL  
Entity: the past few days, Type: DATE  
Entity: about only one, Type: CARDINAL  
Entity: about four days, Type: DATE  
Entity: the year, Type: DATE  
Entity: the afternoon, Type: TIME  
Entity: Andalusia, Type: GPE  
Entity: Moorish, Type: NORP  
Entity: the two hours, Type: TIME  
Entity: each day, Type: DATE  
Entity: Andalusian, Type: NORP  
Entity: one, Type: CARDINAL  
Entity: Moorish, Type: NORP
```

## Evaluation using F1 Score

To calculate F1 score, we need to calculate precision and recall. For this, we first need to manually label the paragraphs we have chosen. The paragraphs are manually labelled using the entities used in tagging.

```
In [17]: # manually labelling the entities of all paras
para_1HumanAnnotation = [("Santiago", "PERSON"), ("the night", "TIME"), ("next day", "DATE"), ("a week ago", "DATE")]

para_2HumanAnnotation = [("second", "ORDINAL"), ("the night", "TIME"), ("a few hours", "TIME"), ("Spain", "GPE"), ("until nightfa

para_3HumanAnnotation = [("Tarifa", "PERSON"), ("midday", "TIME")]
```

```
In [20]: calculatesscore(para_1HumanAnnotation, para_1NER)
```

```
Precision: 1.00  
Recall: 1.00  
F1 Score: 1.00
```

```
In [21]: calculatesscore(para_2HumanAnnotation,para_2NER)
```

```
Precision: 1.00  
Recall: 0.60  
F1 Score: 0.75
```

```
In [22]: calculatesscore(para_3HumanAnnotation,para_3NER)
```

```
Precision: 1.00  
Recall: 0.50  
F1 Score: 0.67
```

### Paragraph-1:

1. Precision = 1.00
2. Recall = 1.00
3. F1 score = 1.00

### Paragraph-2:

1. Precision = 1.00
2. Recall = 0.60
3. F1 score = 0.75

### Paragraph-3:

1. Precision = 1.00
2. Recall = 0.50
3. F1 score = 0.67

## Visualization :-

```
In [19]: displacy.render(para_1NER, jupyter = True, style = 'ent')
```

The boys name was **Santiago PERSON** dusk was falling as the boy arrived with his herd at an abandoned church The roof had fallen in long ago and an enormous sycamore had grown on the spot where the sacristy had once stood He decided to spend **the night TIME** there He saw to it that all the sheep entered through the ruined gate and then laid some planks across it to prevent the flock from wandering away during **the night TIME** There were no wolves in the region but once an animal had strayed during **the night TIME** and the boy had had to spend the entire **next day DATE** searching for it He swept the floor with his jacket and lay down using the book he had just finished reading as a pillow He told himself that he would have to start reading thicker books they lasted longer and made more comfortable pillows It was still dark when he awoke and looking up he could see the stars through the halfdestroyed roof I wanted to sleep a little longer he thought He had had the same dream that night as **a week ago DATE** and once again he had awakened before it ended

## TF-IDF Vector

TF-IDF (Term Frequency-Inverse Document Frequency) is a numerical statistic in Natural Language Processing that reflects the importance of a word in a document relative to a collection of documents (corpus). It is commonly used for text mining and information retrieval. The formula for TF-IDF is given by:

**TF = (Occurrence of a word in a document) / (Number of words in the document)**

**IDF = Log [ (Total number of documents) / (Number of documents containing the word + 1) ]**

**TF-IDF = TF \* IDF**

TF-IDF is finally used for checking similarity between chapters of the book. The following is process for calculating TF-IDF:

```
In [24]: ch_1 = preprocessed_text[0:49481]
ch_2 = preprocessed_text[49481:98960]
ch_3 = preprocessed_text[98960:148441]
ch_4 = preprocessed_text[148441:197920]
chaptersBook = [ch_1, ch_2, ch_3, ch_4]
```

```
In [23]: def idf(chapters):
# Initialize TF-IDF vectorizer
vectorizer = TfidfVectorizer()


# Compute TF-IDF vectors for chapters
tfidf_matrix = vectorizer.fit_transform(chapters)

# Compute similarity scores between chapters using cosine similarity
similarity_matrix = cosine_similarity(tfidf_matrix)

# Create a DataFrame to visualize similarity scores
chapter_names = [f"Chapter {i+1}" for i in range(len(chapters))]
similarity_df = pd.DataFrame(similarity_matrix, columns = chapter_names, index = chapter_names)

plt.figure(figsize=(24, 18))
sns.heatmap(similarity_df, annot=True,)
plt.title("Chapter Similarity using TF-IDF Vectors")
plt.xlabel("Chapter")
plt.ylabel("Chapter")
plt.show()
```

```
In [25]: idf(chaptersBook)
```



The table showcases the similarity of each chapter with every other chapter using the TF-IDF vectors generated earlier.

## Inferences

The following observations were made after using the NLP frameworks on the document 'The Alchemist':

1. From the TF-IDF similarity matrix, we can see that the similarity between same chapters is always 1 which is obvious.
2. The TF-IDF similarity matrix is a symmetric positive definite matrix.